

Certified lower bounds for correlation clustering on large sparse graphs

Quang-Vinh Dang^{1*}

^{1*}British University Vietnam, Hung Yen, Vietnam.

Corresponding author(s). E-mail(s): vinh.dq4@buv.edu.vn;

Abstract

Correlation clustering (equivalently, cluster editing) has strong approximation algorithms and strong heuristics, but on large graphs nobody can say how far a given clustering is from optimal: exact solvers stop at a few hundred vertices, and the lower bounds that scale certify only a factor of about two. We compute certified lower bounds on sparse graphs with up to 10^6 vertices. Since optimal clusters have diameter at most two, the metric LP and its strengthening by star and local subgraph inequalities can be restricted to the $O(\sum_v \deg(v)^2)$ pairs at distance at most two without losing validity or LP-rounding guarantees. Dual block-coordinate ascent, warm-started by a sparse star packing, turns this relaxation into an anytime bound that is valid after every step. Its certificates are verified against the raw instance file, in exact arithmetic, by a short checker that shares no code with the solver; the inequality it evaluates is proved in Lean 4. The optimality gap of any clustering decomposes exactly into non-negative local terms, which yields local optimality certificates and guides a large-neighbourhood search. On the PACE 2021 exact track the checked bound proves optimality on 106 of 173 solved instances and improves the published lower bounds of two open instances. On 18 SNAP graphs, 7 of them held out during development, the better of this bound and the star packing alone certifies the best known solutions within 0.25% to 23.5%, where scalable baselines certify factors of 1.17 to 1.90.

Keywords: correlation clustering, cluster editing, lower bounds, linear programming, dual ascent, certificates, verified computation

MSC Classification: 90C27 , 90C05 , 90C57 , 05C85 , 68W25

1 Introduction

In *correlation clustering* [1] the input is a complete graph whose pairs are labelled *similar* (+) or *dissimilar* (−), and the task is to partition the vertices so that the number of disagreements, + pairs between clusters plus − pairs inside clusters, is minimum. The number of clusters is free. Equivalently, given the positive graph $G^+ = (V, E^+)$, one adds or deletes the fewest edges that turn G^+ into a disjoint union of cliques (*cluster editing*). The problem is APX-hard [2]. It arises in entity resolution, community detection and the consolidation of inconsistent pairwise information.

Two lines of work have progressed separately. Approximation ratios for complete graphs have fallen from 3 for Pivot [3] through 2.06 by rounding the metric LP [4] to 1.485 with the cluster LP [5] and below [6]. In practice, engineered heuristics such as KaPoCE [7], the winner of the PACE 2021 challenge [8], and multilevel memetic solvers [9] find very good solutions on graphs with millions of vertices. What is missing on such graphs is a *certificate*: a lower bound on the optimum, checkable independently of the solver, that tells how far a given clustering can be from optimal. Exact branch-and-bound solves instances with a few hundred vertices [10], and its strong packing bounds are implemented on a dense $n \times n$ matrix. Solvers for the metric LP reach about 10^4 vertices [11–13]. The combinatorial bounds that do scale, packings of bad triangles and the LP relaxations based on strong triadic closure [14], certify solutions only within a factor of about two.

This paper is about the lower-bound side. We compute lower bounds for correlation clustering on sparse graphs with up to 1,088,092 vertices and 1,541,898 edges, export them as certificates that an independent program checks in exact arithmetic against the raw instance file, and use them to certify the solutions of our own and of other solvers. The primal methods are secondary; they are included because a bound is only as useful as the solutions it is compared with, and because the certificate yields a decomposition of the optimality gap that guides a local search.

Contributions.

1. **Distance-two relaxation** (Section 3.1). Optimal clusters have diameter at most two in G^+ , so every pair at distance at least three is separated in every optimal clustering. We show that the metric LP restricted to the pairs at distance at most two, with the triangle inequalities through distance-three pairs kept as covering rows, is a relaxation, and that each of its feasible points extends to a feasible point of the full metric LP with the same cost (Theorem 3). The relaxation has $O(\sum_v \deg(v)^2)$ variables instead of $\binom{n}{2}$, and LP rounding keeps its guarantees. We add star and local subgraph inequalities, which are valid for every clustering that separates the distance-three pairs. This observation is simple; its value is that it makes the metric LP and its strengthenings usable on graphs with 10^5 to 10^6 vertices.
2. **Anytime block-dual bound** (Section 3.2). We compute a lower bound by dual block-coordinate ascent: each step solves the LP of one vertex block exactly by cutting planes, with costs equal to the current reduced costs. The bound is valid after every step and never decreases, and its state is one reduced cost per pair plus the rows with positive multipliers (Theorem 5). Any packing of induced stars or




Fig. 1 Overview. The dual side (orange) builds the distance-two support and a lower bound by block dual ascent from a sparse star packing; the certificate is re-checked by an independent program in exact arithmetic against the raw instance. The primal side (blue) computes clusterings: Pivot and iterated flipping give a start, and either the gap-guided LNS of CertiFlip or the memetic search PXM improves it. The gap decomposition splits $\text{cost}(\mathcal{C}) - \text{LB}$ into non-negative local terms and tells the LNS where to search.

bad triangles is a feasible start. We compute one on the sparse support by a ruin-and-recreate local search. This sparse star packing is the scalable counterpart of the root bound of the KaPoCE branch-and-bound. On large sparse graphs a long packing alone is stronger than a short packing followed by block LPs in the same budget; the block LPs pay off on small and dense instances (Sections 7.2 and 7.3).

3. **Exactly checked certificates** (Section 4). A certificate lists the rows with positive multipliers. A checker of about 505 lines that shares no code with the solver parses the raw instance file, verifies that it is the instance the certificate was computed for (SHA-256 of the file and of the canonical edge list), checks the validity of every row by enumeration or by a closed form, and evaluates the bound in integer arithmetic. The statement that the checker evaluates is proved in Lean 4 (Theorem 5(a)). We archive every certificate behind every reported bound.
4. **Gap decomposition and CertiFlip** (Section 5). For any clustering, the difference between its cost and the dual bound splits exactly into non-negative terms on pairs and rows (Theorem 7). Re-optimising a vertex set can gain at most the gap mass that touches it, which certifies local optimality and selects neighbourhoods for a large-neighbourhood search with exact sub-MIPs. CertiFlip combines this with a flipping local search after Cohen-Addad et al. [15] and returns a checked bound with every solution.
5. **Computational study** (Section 7). On the 173 solved instances of the PACE 2021 exact track, the checked bound proves optimality on 106 and the root star packing of the KaPoCE branch-and-bound on 79; averaged over instances the latter is stronger (0.9948 against 0.9874 of the optimum), and the comparison depends on density. Our bound improves the published lower bounds of two open instances. On 18 SNAP graphs, 7 of them from a held-out set fixed in advance, the better of the

block-dual bound and the star packing alone certifies the best known solutions within 0.25% to 23.5%, where scalable baselines certify factors of 1.17 to 1.90. We compare with the sparse star packing alone and with the metric LP on the support with triangle or star rows, so that the contribution of each part is visible.

A second primal solver, PXMem (Section 6), recombines two clusterings exactly block by block of their overlap. We use it to measure how tight the certificates are against strong solutions, and compare it with KaPoCE under a protocol fixed in advance (Section 7.4). The approximation guarantees we state for CertiFlip and PXMem, expected cost at most 3 OPT, are inherited from the Pivot start and serve only as a safeguard; the quality of a solution is measured by its certificate.

Relation to certified computation.

Certifying algorithms [16] return a witness with every output. Safe bounds from floating-point LP solutions [17, 18], checkable certificates for mixed-integer programming [19] and proof logging in SAT and pseudo-Boolean solving [20, 21] pursue the same goal. Our certificates are simpler than general MIP certificates because every row is local: its validity is decided by enumerating the partitions of at most nine vertices or by a closed form.

Development and evaluation data.

The methods were developed on fifteen SNAP graphs and the PACE 2021 exact track. Before the final experiments we fixed a held-out set of twelve further SNAP graphs and the analysis plan (Section 7.1); the PACE heuristic-track instances that coincide with development graphs are reported separately. Appendix B records which data informed which design decision.

2 Preliminaries and related work

Let $G^+ = (V, E^+)$ with $n = |V|$ and $m = |E^+|$; every pair not in E^+ is negative. For a clustering $\mathcal{C}$ let $x_{uv}^{\mathcal{C}} = 1$ if u, v are in different clusters and 0 otherwise. Then

$$\text{cost}(\mathcal{C}) = \sum_{uv \in E^+} x_{uv}^{\mathcal{C}} + \sum_{uv \notin E^+} (1 - x_{uv}^{\mathcal{C}}).$$

$N(v)$ denotes the positive neighbourhood of v and $d(u, v)$ the distance in G^+ . The *metric LP* [2, 3] relaxes x to $[0, 1]^{\binom{V}{2}}$ subject to the triangle inequalities $x_{uv} \leq x_{uw} + x_{wv}$ for all triples. A *bad triangle* is a triple u, w, v with $uw, wv \in E^+$ and $uv \notin E^+$; every clustering makes at least one mistake on it. An *induced star* with centre v and $k \geq 2$ leaves is a set of k pairwise non-adjacent neighbours of v ; every clustering makes at least $k - 1$ mistakes on its $k + \binom{k}{2}$ pairs.

Approximation.

Pivot repeatedly picks a uniformly random unclustered vertex and clusters it with its unclustered positive neighbours; its expected cost is at most three times the fractional

packing LP of bad triangles [3]. LP-based pivoting joins v to the pivot u with probability $1 - f(x_{uv})$; $f(x) = x$ gives 2.5 [3], and the rounding functions of Chawla et al. [4] give 2.06. Both guarantees hold against the value of *any* feasible metric x , a fact we use below. Further improvements use Sherali–Adams hierarchies [22], preclustering [23] and the cluster LP [5, 6, 24]; the best combinatorial ratio, $2 - 2/13$, comes from a local search with flips [15]. For parallel, streaming, dynamic and weighted variants see [25–27].

Lower bounds.

The classical scalable bounds are packings of bad triangles, greedy or fractional; the latter approximates the covering LP over bad triangles and can be computed by the multiplicative-weights scheme of Garg and Könemann [28, 29]. Veldt [14] derives LP relaxations from strong triadic closure that are supported on the open wedges of G^+ , like our relaxation, and yield fast 2-approximations and bounds within a factor of about two; our relaxation keeps the full metric LP on these pairs and adds stronger rows. The metric LP itself has been solved on graphs with about 10^4 vertices by projection methods [11, 12] and by the Project and Forget scheme [13], which also keeps only active triangle rows. Exact branch-and-bound for cluster editing [10, 30] uses packings of stars and of induced paths on three vertices as lower bounds; Bläsius et al. [10] mark the pairs at distance at least three as forbidden, and Bastos et al. [31] observe that they are separated in every optimal solution. For the clique partitioning problem, the weighted maximisation form, facets and cutting planes go back to Grötschel and Wakabayashi [32, 33]; later work removes redundant triangle rows [34–36], adds chorded-cycle facets [37], uses Lagrangian relaxation with pegging tests [38], and tightens bounds with constraints on subnetworks [39]. Exact approaches via weighted MaxSAT [40] and kernelisation [41, 42] solve smaller instances to optimality.

For the multicut problem on sparse graphs, where missing pairs carry no cost, dual decomposition and message passing give fast lower bounds [43–45], and persistency criteria fix variables that take the same value in every optimal solution [46, 47]. Our block-dual ascent is closest in spirit to block-coordinate ascent for dense graphical models [48]: each block is solved to optimality with the other multipliers fixed. These multicut methods do not apply directly to correlation clustering on complete graphs, where the $\binom{n}{2} - m$ negative pairs all carry cost; the distance-two support is what reduces the complete instance to a sparse one without changing the optimum.

Heuristics.

Local search with vertex moves [49], Louvain and Leiden with the constant Potts model at resolution $1/2$, which optimise the same objective [50–52], heuristics for clique partitioning [53–55], local search for multicut [56] and fusion moves [57] are the common primal methods. KaPoCE combines reductions, local search and an evolutionary phase [7]. SCC [9] is a multilevel memetic solver for correlation clustering on sparse signed graphs, in which pairs without an edge carry no cost; its objective therefore differs from the complete-graph objective studied here (Section 7.1). Twin (critical-clique) reductions and their weighted variants are standard in parameterised algorithms and practical tools [41, 58–63].




Fig. 2 The distance-two support on a 40-vertex neighbourhood of **ca-GrQc**. Left: the positive graph. Right: the support P adds the non-adjacent pairs with a common neighbour (orange); every other pair is far, separated in every optimal clustering (Lemma 1), and carries no variable.

3 Lower bounds on the distance-two support

3.1 Distance-two relaxations

Lemma 1 *Let $\mathcal{C}$ be an optimal clustering, $S \in \mathcal{C}$ and $v \in S$. Then $|N(v) \cap S| \geq (|S| - 1)/2$. Consequently any two vertices of S are adjacent or have a common neighbour in S . In particular every pair at distance at least three is separated.*

Proof Moving v to a singleton changes the cost by $|N(v) \cap S| - (|S| - 1 - |N(v) \cap S|)$, which must be non-negative. For non-adjacent $u, v \in S$ the sets $N(u) \cap S$ and $N(v) \cap S$ lie in $S \setminus \{u, v\}$, which has $|S| - 2$ elements, and have total size at least $|S| - 1$, so they intersect. $\square$

Lemma 1 is folklore; its consequence for distance-three pairs is noted by Bastos et al. [31] and used to forbid pairs in the branch-and-bound of Bläsius et al. [10]. What follows uses it to restrict the LP rather than the search.

Let $N_2 = \{uv \notin E^+ : N(u) \cap N(v) \neq \emptyset\}$, $P = E^+ \cup N_2$ and $F = \binom{V}{2} \setminus P$ (the *far* pairs). Note $|P| \leq m + \sum_w \binom{\deg(w)}{2}$. Figure 2 shows the support on a small neighbourhood of a collaboration network: already there, 587 of the 780 pairs (75%) are far and disappear from the relaxation.

Definition 2 (distance-two relaxation) LP_P has a variable $x_e \in [0, 1]$ for every $e \in P$ and minimises $\sum_{e \in E^+} x_e + \sum_{e \in N_2} (1 - x_e)$ subject to

$$x_{uv} \leq x_{uw} + x_{wv} \quad \text{for all } u, v, w \text{ with } uv, uw, wv \in P, \quad (1)$$

$$x_{uw} + x_{wv} \geq 1 \quad \text{for all } u, v, w \text{ with } uw, wv \in P, uv \in F. \quad (2)$$

Theorem 3 Let LP_{std} denote the value of the metric LP.

- (a) $\text{LP}_{\text{std}} \leq \text{LP}_P \leq \text{OPT}$.
- (b) If x is feasible for LP_P , then $\bar{x}$, defined as x on P and 1 on F , is feasible for the metric LP and has the same cost.
- (c) Let $\mathcal{A}$ be a randomised rounding algorithm with $\mathbb{E}[\text{cost}(\mathcal{A}(\bar{x}))] \leq \alpha \text{cost}(\bar{x})$ for every feasible metric $\bar{x}$. Applied to an optimal solution of LP_P it is an α -approximation. This holds in particular for $\alpha = 2.5$ (ACN) and $\alpha = 2.06$ (CMSY). With pivot-based rounding, only pairs in P can be joined, so the rounding takes $O(n + |P|)$ time.

Proof (b) Consider a triple u, v, w and the number of far pairs among uv, uw, vw . With none, all three triangle inequalities are rows (1). With exactly one, say uv , the inequality $\bar{x}_{uv} = 1 \leq x_{uw} + x_{vw}$ is row (2). The other two inequalities, such as $x_{uw} \leq 1 + x_{vw}$, hold because $x \leq 1$. With two or three far pairs, each inequality has a right-hand side of at least 1. Far pairs are negative and $\bar{x} = 1$ on them, so they contribute 0 to the cost. (a) The first inequality follows from (b). For the second, by Lemma 1 an optimal clustering separates all far pairs. Its cut vector restricted to P therefore satisfies (1) and (2), because it is a metric with $x = 1$ on F , and it has cost OPT . (c) follows from (a) and (b): the expected cost is at most $\alpha \text{cost}(\bar{x}) = \alpha \text{LP}_P \leq \alpha \text{OPT}$. Pivot-based rounding joins v to the pivot u with probability $1 - f(x_{uv})$, and $f(1) = 1$ for both rounding functions. $\square$

The rows (1)–(2) are generated lazily. A violated row needs $x_{uw} + x_{vw} < 1$, so separation around a vertex w only has to look at pairs of “close” neighbours in P , which keeps it fast in practice.

Stronger inequalities.

The integrality gap of the metric LP is at least 2 (stars attain it asymptotically) and at most 2.06 [4]. We add two families that are valid for every clustering separating all far pairs, and hence for every optimal clustering.

Proposition 4 Let $\mathcal{C}$ be a clustering in which no far pair shares a cluster.

- (i) (star inequalities; the case $|S| = 1$ of the 2-partition inequalities of Grötschel and Wakabayashi [32]) For every v and $T \subseteq V \setminus \{v\}$, with $y = 1 - x^{\mathcal{C}}$,

$$\sum_{t \in T} y_{vt} - \sum_{\{t, t'\} \subseteq T, tt' \in P} y_{tt'} \leq 1.$$

- (ii) (local subgraph inequalities) For every $H \subseteq V$,

$$\sum_{e \in E^+(H)} x_e^{\mathcal{C}} + \sum_{e \in N_2(H)} (1 - x_e^{\mathcal{C}}) \geq \text{OPT}(H),$$

where $\text{OPT}(H)$ is the optimum of the instance induced by H and $E^+(H)$, $N_2(H)$ are the pairs of E^+ and N_2 inside H .

Proof (i) If the cluster of v contains k vertices of T , the first sum is k and the second is at least $\binom{k}{2}$. Pairs inside T that are not in P are far, so they are separated and contribute 0 to the second sum. Finally $k - \binom{k}{2} \leq 1$. (ii) Restricted to H , $\mathcal{C}$ is a clustering of H . Its mistakes inside H number at least $\text{OPT}(H)$. The far pairs inside H are separated and cost nothing, so all its mistakes are counted on the left-hand side. $\square$

The local subgraph inequalities are local cuts in the sense of Applegate et al. [64]: valid inequalities obtained by solving a small subproblem exactly. Bad triangles ($\text{OPT} = 1$) and induced stars with k independent leaves ($\text{OPT} = k - 1$) are local subgraph inequalities. So are chordless cycles C_k , for which $\text{OPT}(C_k) = \lceil k/2 \rceil$ while the metric LP only certifies $k/2$ (put $x = 1/2$ on the cycle edges). We separate stars greedily around each vertex. For (ii) we grow H ($|H| \leq 8$) by breadth-first search along fractional pairs and compute $\text{OPT}(H)$ by enumerating the at most 4140 set partitions of H .

3.2 Anytime dual bounds

Write the relaxation as $\min\{K + c^\top x : Ax \geq b, 0 \leq x \leq 1\}$ over $x \in \mathbb{R}^P$, with $c_e = 1$ on E^+ , $c_e = -1$ on N_2 and $K = |N_2|$. The rows of (A, b) are any valid inequalities from Section 3.1.

Theorem 5 *For every $y \geq 0$ let $r = c - A^\top y$ and*

$$\text{LB}(y) = K + b^\top y + \sum_{e \in P} \min\{0, r_e\}. \quad (3)$$

- (a) $\text{LB}(y) \leq \text{OPT}$.
- (b) Let $B \subseteq V$, let R_B be a set of rows supported on pairs with both endpoints in B , and let $\bar{r}$ be the reduced costs after removing the contribution of the current multipliers on R_B . Replace these multipliers by an optimal dual solution of $\min\{\bar{r}^\top x_B : A_B x_B \geq b_B, 0 \leq x_B \leq 1\}$. This does not decrease LB .
- (c) If $B = V$ and R_B contains all rows of LP_P , one block step gives $\text{LB} = \text{LP}_P$.

Proof (a) Let $\mathcal{C}$ be optimal. By Lemma 1 it separates far pairs, so $\text{OPT} = K + c^\top x^{\mathcal{C}}$ with $x^{\mathcal{C}} \in \{0, 1\}^P$ satisfying $Ax^{\mathcal{C}} \geq b$. Then $\text{OPT} = K + r^\top x^{\mathcal{C}} + y^\top Ax^{\mathcal{C}} \geq K + \sum_e \min\{0, r_e\} + y^\top b$. (b) By LP duality the block LP equals $\max_{y_B \geq 0} b_B^\top y_B + \sum_{e \in P_B} \min\{0, (\bar{r} - A_B^\top y_B)_e\}$. The terms of (3) outside the block do not change, and the old multipliers are feasible for this maximisation. (c) is LP duality for LP_P with the box constraints. $\square$

Part (b) needs an optimal block dual. When a block LP stops at its time limit, its duals are only feasible; the bound computed with them is still valid by part (a), but it may be smaller than before. The implementation therefore evaluates the new bound and keeps a block step only if the bound does not decrease, so the bound is monotone in every case.

Warm start by star packings.

Every pair-disjoint family of induced stars (a centre v with $k \geq 2$ pairwise non-adjacent positive neighbours) is a feasible dual solution: put $y = 1$ on the corresponding local subgraph inequalities of Proposition 4(ii), which in x -form read $\sum_t x_{vt} - \sum_{tt'} x_{tt'} \geq (k-1) - \binom{k}{2}$. Then $\text{LB}(y)$ equals the packing value $\sum (k-1)$. Indeed, a centre pair has $c_e = 1$ and coefficient 1, and a leaf pair has $c_e = -1$ and coefficient -1 , so every packed pair has reduced cost 0. Unpacked pairs keep $r_e = c_e$. Hence $\text{LB} = K + \sum_S ((k_S - 1) - \binom{k_S}{2}) - |N_2 \setminus \text{packed}|$, and $K - |N_2 \setminus \text{packed}| = \sum_S \binom{k_S}{2}$. Fractional packings of bad triangles, as computed by the Garg–Könemann scheme [28, 29], are installed in the same way on the rows $x_{uw} + x_{wv} - x_{uv} \geq 0$. Packings are the bounds used inside the KaPoCE branch-and-bound algorithm [10]. Its implementation stores the instance as a dense matrix, so it applies only to graphs with at most about 10^4 vertices. We build a greedy packing on the sparse support, processing centres by decreasing degree, and improve it by *ruin and recreate*. (On dense instances large stars waste pairs, since a star with k leaves uses $k + \binom{k}{2}$ pairs for a value of $k-1$. There we also try a packing that first packs bad triangles maximally and then extends stars on free pairs, as well as the fractional triangle packing, and install the best of the candidates.) Each step picks a random vertex, removes all stars centred in a random connected region of about 12 vertices around it, and repacks the freed pairs greedily in a random order. A step is kept iff the value does not decrease. The packing is installed as the initial multipliers. Subsequent block steps release the packed rows that lie inside a block and re-optimize them fractionally together with newly separated rows. By Theorem 5(b), the final bound is never below the packing value.

The bound (3) needs only the vector r and the scalar $b^\top y$. We store the rows with positive multipliers so that a later block containing them can re-optimize them. Blocks of the same partition are independent. In a *sweep* we compute a partition into parts of about 2500 vertices with METIS [65] on a randomly relabelled copy of G^+ , so that successive sweeps cut different pairs. Each part is solved with the dual simplex method of HiGHS [66], with lazily separated triangle, star and local subgraph rows. After the sweeps stall we switch to *gap-guided blocks*, grown around the vertices with the largest share of the gap of Section 5. Because (3) is valid for any $y \geq 0$, the bound remains correct even when a block LP hits its time limit. We only clamp the dual values to be non-negative, and reject a block step that would lower the bound.

Floating point.

The block LPs are solved in double precision, so the multipliers y are inexact and the rows they carry are generated by floating-point separation. Neither affects validity: (3) holds for every $y \geq 0$ and every set of valid rows. The value the solver reports is nevertheless not the value we publish. Section 4 recomputes it from the exported rows and multipliers in exact arithmetic, after checking every row, and all bounds in Section 7 are these recomputed values.

4 Certificates and their independent check

Every lower bound reported in Section 7 is the output of a separate checker, not of the solver. This section specifies what a certificate contains, what the checker verifies, and what has to be trusted.

Contents.

A certificate is a compressed NumPy archive with

- the rows with positive multiplier, in compressed sparse row form: for row i the pairs $p = \{u, v\}$ of its support, integer coefficients a_{ip} and an integer right-hand side b_i , in the x -form $a_i^\top x \geq b_i$;
- the multipliers $y_i \geq 0$ as double-precision numbers;
- the identity of the instance: the name and relative path of the raw file as distributed by PACE or SNAP, its SHA-256, its format, the number of vertices n and edges m , and the SHA-256 of the canonical edge list, that is, of n followed by the lexicographically sorted pairs (u, v) with $u < v$ as 64-bit integers.

The vertex numbering is fixed by the instance definition: a PACE file numbers the vertices $1, \dots, n$; in a SNAP edge list the vertices are the identifiers that occur in a kept edge, numbered in increasing order, a signed list keeps only its positive rows, and self-loops and duplicate edges are dropped. The clustering whose cost is reported is archived next to the certificate, one label per vertex in the same numbering.

The check.

The checker (`experiments/check_certificate.py`, 505 lines of Python with Numba kernels) imports nothing from the solver package. Given the raw file and the certificate it

1. parses the raw file with its own reader and rejects the certificate unless n , m , the file hash and the edge-list hash agree with the recorded ones;
2. checks that every pair of every row is in P , i.e. is an edge or has a common neighbour, by intersecting sorted adjacency lists;
3. checks that every row is valid for every clustering that separates the far pairs. A row on at most nine vertices is checked by enumerating all partitions of its vertex set (at most 21,147, as restricted growth strings) that separate its far pairs and taking the minimum of the integer left-hand side. A larger row is accepted only if it is a star row: its positive entries are the pairs vt , $t \in T$, of one centre v , and its negative entries are pairs R inside T . Its minimum over far-separating clusterings is at least $|T| - |R| - M$ with $M = 1$ if R contains every non-far pair inside T (then a far-free subset of T is a clique of R) and M the independence number of (T, R) otherwise, computed exactly;
4. rounds each y_i down to a multiple of 2^{-30} and evaluates

$$\text{LB} = \sum_i b_i y_i + \sum_{p \in P} \left(\min\{0, c_p - \sum_i a_{ip} y_i\} - \min\{0, c_p\} \right) \quad (4)$$

in integer arithmetic (64-bit accumulation under a checked overflow bound, then arbitrary-precision integers), and returns $\lceil \text{LB} \rceil$;

5. optionally evaluates the cost of the archived clustering on the parsed instance, also in integers.

Rounding down keeps $y \geq 0$, so the rounded multipliers are again a dual certificate; nothing else in (4) is inexact. The sum runs over the pairs that occur in some row, since the other terms vanish. Expression (4) is (3) with $K = |N_2| = -\sum_{p \in P} \min\{0, c_p\}$ absorbed.

Proposition 6 *If the checker accepts a certificate for an instance, then the value it returns is at most OPT of that instance.*

Proof By steps 2 and 3, every row is supported on P and satisfied by the cut vector of every clustering that separates far pairs. By Lemma 1 this includes every optimal clustering, and Theorem 5(a) with the rounded multipliers gives $\text{LB} \leq \text{OPT}$. OPT is an integer. $\square$

The inequality behind the proof, that (4) is at most the cost of every far-separating clustering for any rows valid for such clusterings, is proved in Lean 4 (`far_dual_bound`, `far_dual_le_opt`; Appendix A), as is the fact that optimal clusterings separate far pairs. What remains to be trusted is the checker’s reader, the enumeration and the star formula of step 3, the integer evaluation, and the Python and NumPy runtime. The tests of the repository check the reader against the solver’s loader (same n , m and edge hash on all instances we use), the star formula against the enumeration on random star rows with at most nine vertices, the integer cost of a clustering against the solver’s, and the rejection of certificates with a perturbed right-hand side, a negative multiplier, a pair outside P , or a different instance.

Cost of checking.

Checking is much cheaper than computing the bound: the largest certificate we report, on `roadNet-PA`, has 610,758 rows and is checked in 8 s (Section 7.3). One command re-checks every archived certificate against freshly downloaded instance files and writes one line per certificate.

What is not certified.

The certificate proves a lower bound, and the archived clustering an upper bound; together they bound the gap of the reported solution. The approximation guarantees of Theorems 3 and 9 are statements about the algorithms, not about a particular run, and are not checked. Running times are measured, not certified.

5 Gap decomposition and CertiFlip

A dual solution does more than bound the optimum: together with a clustering it says where the clustering can still improve. This section makes that precise and uses it in a primal method that returns a certificate with every solution.

Theorem 7 Let $y \geq 0$, $r = c - A^\top y$, and let $\mathcal{C}$ be a clustering that separates all far pairs. Then

$$\text{cost}(\mathcal{C}) - \text{LB}(y) = \sum_{e \in P} \underbrace{(r_e x_e^{\mathcal{C}} - \min\{0, r_e\})}_{g_e \geq 0} + \sum_i y_i \underbrace{(a_i^\top x^{\mathcal{C}} - b_i)}_{s_i \geq 0}. \quad (5)$$

Proof $\text{cost}(\mathcal{C}) = K + c^\top x^{\mathcal{C}} = K + r^\top x^{\mathcal{C}} + y^\top A x^{\mathcal{C}}$. Subtract (3). Each $g_e \geq 0$ because $x_e \in \{0, 1\}$, and each $s_i \geq 0$ by the validity of row i . $\square$

A pair term vanishes iff the clustering agrees with the sign of the reduced cost. A row term vanishes iff the row is tight or has multiplier 0. Assigning each term to the endpoints of its pairs gives a *gap map* over the vertices.

Corollary 8 (local certificate) For $B \subseteq V$ let $\Gamma_B(\mathcal{C})$ be the sum of the terms of (5) whose pair, or one of whose row pairs, has an endpoint in B . Let $\mathcal{C}'$ be any clustering that separates far pairs and coincides with $\mathcal{C}$ on all pairs with both endpoints outside B . Then $\text{cost}(\mathcal{C}) - \text{cost}(\mathcal{C}') \leq \Gamma_B(\mathcal{C})$. In particular, if $\Gamma_B(\mathcal{C}) < 1$ then no re-clustering of B improves $\mathcal{C}$.

Proof Apply (5) to both clusterings. Terms not touching B coincide, and the terms of $\mathcal{C}'$ are non-negative. Costs are integers. $\square$

Gap-guided LNS.

Large-neighbourhood search [67], here with exact sub-MIPs in the spirit of RINS and local branching [68, 69], fixes the clustering outside a vertex set B ($|B| \leq 60$) and re-optimises B exactly. Let $\mathcal{K}$ be the outside clusters with a positive edge into B . Re-clusterings of B correspond to clusterings of B together with one super node per $K \in \mathcal{K}$, where super nodes must stay apart. The cost is $\sum_{uv} [uv \in E^+] x_{uv} + [uv \notin E^+] (1 - x_{uv})$ for $u, v \in B$, plus $w_{vK} x_{vK} + (|K| - w_{vK})(1 - x_{vK})$ for each $v \in B$, where w_{vK} counts the positive edges between v and K . We solve this MIP with HiGHS and add triangle inequalities lazily until the solution is transitive. To keep Corollary 8 applicable, v may join K only if it is within distance two of every member of K , and far pairs inside B are fixed apart. Neighbourhoods are grown around high-gap vertices, either by breadth-first search or as a union of whole clusters. Neighbourhoods with $\Gamma_B < 1$ are skipped because they are certified optimal. Theorem 7 and Corollary 8 need a clustering that separates far pairs, which the output of the flipping search need not do. Before the search, and after every improvement, we therefore move to a singleton every vertex v with $|N(v) \cap S| < (|S| - 1)/2$ in its cluster S , until there is none. Each such move lowers the cost, and by the argument of Lemma 1 the result separates all far pairs. The certificate $\Gamma_B < 1$ covers re-clusterings of B that also separate far pairs; by Lemma 1 these include every optimal one.

Figure 3 shows the gap map at work on the neighbourhood of Figure 2. For a Pivot clustering the gap of 10 is concentrated on one cluster; the LNS re-optimises the blocks around the high-gap vertices and reaches a clustering whose cost equals the bound, which certifies it optimal.




Fig. 3 Gap decomposition (Theorem 7) on the neighbourhood of Figure 2. Each vertex is coloured by its share of $\text{cost}(\mathcal{C}) - \text{LB}$ (pair terms split between the endpoints, row terms between the vertices of the row); shaded regions are clusters, dark edges lie inside a cluster. Left: a Pivot clustering. Right: after gap-guided LNS the gap is zero, so the clustering is optimal. Computed by the released code.

Algorithm 1 CertiFlip(G^+, T)

- 1: $\mathcal{C} \leftarrow \text{PIVOT}(G^+)$ with a uniformly random order
 - 2: $\mathcal{C} \leftarrow \text{ITERATEDFLIP}(\mathcal{C})$ ▷ insertion local search with flips
 - 3: build the distance-two support P
 - 4: $y \leftarrow$ star packing with ruin-and-recreate local search ▷ feasible dual
 - 5: $y \leftarrow \text{BLOCKASCENT}(P, y, \mathcal{C})$ ▷ METIS sweeps, then gap-guided blocks
 - 6: **if** one block covered V and separation converged **then**
 - 7: $\mathcal{C} \leftarrow$ best of $\mathcal{C}$ and $\text{ITERATEDFLIP}(\text{CMSY-ROUND}(x^*))$ over several roundings
 - 8: **end if**
 - 9: $\mathcal{C} \leftarrow \text{GAPLNS}(\mathcal{C}, y)$ until time T
 - 10: **return** $\mathcal{C}$ and $\lceil \text{LB}(y) \rceil$ ▷ lines 4–5 share half of the budget T
-

5.1 CertiFlip

Insertion local search.

For every root r we build one candidate insertion $S(r)$ by *cleaning* $N[r]$. Starting from $S = N[r]$, we repeat a few sweeps over $v \in N[r] \setminus \{r\}$: we put v in S iff joining S costs v less than staying in the remainder of its current cluster. Both costs are computed exactly from counts of weighted positive neighbours. We then compute the exact change of the objective in $O(\sum_{v \in S} \deg v)$ time and apply the insertion if it is negative. Rounds of insertions alternate with a multilevel (Louvain-type) vertex-move local search. That search works on the weighted objective, moves super nodes at coarse levels and hence also merges clusters. **ITERATEDFLIP** follows Cohen-Addad et al. [15]. Positive weights take values in $\{1, 1.5, 2\}$ (we double all weights to stay in integers). Each round flips the pairs cut by the previous solution, re-optimises, flips

and re-optimises again, and combines the last three solutions with their 3-way pivot. The best candidate under the original objective is returned.

Theorem 9 *Let $\mathcal{C}$ and LB be the output of Algorithm 1.*

- (a) $\mathbb{E}[\text{cost}(\mathcal{C})] \leq 3 \text{ OPT}$.
- (b) $\text{LB} \leq \text{OPT}$, so $\text{cost}(\mathcal{C})/\lceil \text{LB} \rceil$ is an upper bound on the approximation ratio achieved on the instance.
- (c) If line 7 is executed, $\mathbb{E}[\text{cost}(\mathcal{C})] \leq 2.06 \text{ OPT}$.

Proof Every step after line 1 returns a clustering no worse than its input: the local searches only apply improving moves, ITERATEDFLIP returns the best of its candidates, which include the local optimum reached from its input, and GAPLNS accepts only improvements. (a) then follows from the Pivot guarantee [3]. (b) is Theorem 5(a) and the integrality of OPT. (c) If one block covered V and separation converged, the block LP primal x^* satisfies all rows (1)–(2), so it is feasible for LP_P . Its cost is the value of the block LP. Every row of the block LP is valid for an optimal clustering (Lemma 1 and Proposition 4), whose cut vector is therefore feasible for the block LP; hence the block LP value is at most OPT. By Theorem 3(b,c), $\mathbb{E}[\text{cost}(\text{CMSY-ROUND}(x^*))] \leq 2.06 \text{ OPT}$. The remaining steps are monotone. $\square$

Remark 10 Part (a) holds for any procedure that starts from a Pivot clustering and never makes its solution worse; it is a safeguard, not a distinguishing property. If the insertion local search reached an optimum with respect to *all* cluster insertions, Cohen-Addad et al. [15] would give a factor of 2, and $2 - 2/13$ with enough flipping rounds. Our candidate generation explores a polynomial subfamily of insertions, so we do not claim these factors. The certificate of Theorem 9(b), however, is valid in every case.

Running time.

Pivot takes $O(n + m)$ time. One round of insertions takes $O(\sum_v \sum_{u \in N(v)} \deg u) = O(\sum_u \deg(u)^2)$, the same order as the size of P and the number of bad triangles. The support takes $O(\sum_w \deg(w)^2)$ time and memory. The block LPs dominate the running time, and the time spent on them is set by the budget T .

Budget.

The budget T is checked between steps: before each block LP, each LNS neighbourhood and each flipping round. A single round of the insertion local search, the construction of the support and a sub-MIP are not interrupted, so a run can exceed T . Section 7 reports the measured wall-clock and CPU times of every run, including these overruns, and the time at which the final bound was reached.

6 A primal solver based on exact recombination

A certificate says how far a clustering may be from optimal; how close it actually is depends on the clustering. To test the certificates against strong solutions, and because

CertiFlip spends much of its budget on the bound, we also use a second primal solver, *PXMem*. It rests on two elementary exact facts: twins can be contracted, and two clusterings can be recombined block by block of their overlap without losing anything. Both have close relatives in the literature, which we point out; we claim no novelty for either, only for their combination into a solver. *PXMem* produces no bound of its own; its solutions are certified with the bounds of Section 3.

6.1 Critical-clique contraction

Two vertices $u \neq v$ are *twins* if $uv \in E^+$ and $N(u) \setminus \{v\} = N(v) \setminus \{u\}$, i.e. they have the same closed neighbourhood. The classes of twins are the *critical cliques* of Guo [41]. The following lemma is due to Guo [41] (see also Protti et al. [58]); we include the short proof because the implementation relies on its exact form.

Lemma 11 (41) *Let u, v be twins and $\mathcal{C}$ a clustering that separates them. Moving v into the cluster of u and moving u into the cluster of v change the cost by δ_1 and δ_2 with $\delta_1 + \delta_2 = -2$. Hence every optimal clustering keeps every critical clique inside one cluster.*

Proof Let $A \ni u$ and $B \ni v$ be the two clusters. For a cluster X let $f(X)$ be the number of mistakes on the pairs zw with $w \in V \setminus \{u, v\}$ when a vertex $z \in \{u, v\}$ is placed in X (all other vertices staying where they are). It is the same for u and v because they have the same neighbours in $V \setminus \{u, v\}$. The moves do not change $A \setminus \{u, v\}$ and $B \setminus \{u, v\}$. Besides the cut positive pair uv , the current cost of the pairs involving u or v is $f(A) + f(B)$; after the first move it is $2f(A)$, after the second $2f(B)$. So $\delta_1 + \delta_2 = 2f(A) + 2f(B) - 2(f(A) + f(B) + 1) = -2$ and one of the moves is a strict improvement. $\square$

Contracting every critical clique K to a node of size $s_K = |K|$ gives a weighted instance on the node set U : two nodes $K \neq L$ carry the weight $w_{KL} = s_K s_L$ if K and L are adjacent and 0 otherwise, and the cost of a clustering c of U is

$$\text{cost}_w(c) = \sum_{\{K, L\}: c(K)=c(L)} (s_K s_L - w_{KL}) + \sum_{\{K, L\}: c(K) \neq c(L)} w_{KL}. \quad (6)$$

The pairs inside a critical clique are positive and never cut. Hence the cost of the expanded clustering on G^+ equals $\text{cost}_w(c)$, and by Lemma 11 an optimal clustering of the contracted instance expands to an optimal clustering of G^+ . Merging vertices into weighted nodes in this way is standard in cluster editing [41, 59]. On the collaboration graphs the contraction removes 13%–22% of the vertices. Pivot never separates twins either: whether a vertex joins the cluster of the pivot p depends only on its closed neighbourhood, so each round treats twins alike. A Pivot clustering therefore maps to the contracted instance without changing its cost.

6.2 Partition crossover

Figure 4 illustrates the theorem below on two Pivot clusterings of the neighbourhood of Figure 2.




Fig. 4 Partition crossover (Theorem 12) computed by the released code. (a), (b): two Pivot clusterings of the neighbourhood of Figure 2, both of cost 49; shaded regions are clusters. (c): the connected components of the overlap graph H split the vertices into seven blocks, and on each block the two parents can be compared on their own. (d): taking the better parent on every block gives a clustering of cost 38, better than both parents.

Theorem 12 (partition crossover) *Let the cost of a clustering be any sum of pair terms that depend only on whether the two vertices of the pair share a cluster (this covers cost and the weighted cost (6)). Let a and b be clusterings and let H be the bipartite graph whose nodes are the clusters of a and of b , with an edge between $A \in a$ and $B \in b$ whenever $A \cap B \neq \emptyset$. Each connected component of H covers a vertex set X , the block, which is a union of clusters of a and also a union of clusters of b . For a clustering c let $\text{cost}_X(c)$ count the mistakes on the pairs inside X . Choose, independently in every block, the clusters of b if $\text{cost}_X(b) < \text{cost}_X(a)$ and those of a otherwise. The resulting clustering χ satisfies*

$$\text{cost}(\chi) = \text{cost}(a) - \sum_X \max(0, \text{cost}_X(a) - \text{cost}_X(b)) \leq \min(\text{cost}(a), \text{cost}(b)).$$

Proof A pair u, v in different blocks is separated by a , by b and by χ , because each of their clusters lies inside one block. It therefore contributes the same amount to all three costs. A pair inside a block X is treated by χ exactly as by the parent chosen for X . Summing over the pairs gives the equality. Every term of the sum is non-negative, and exchanging the roles of a and b gives the bound by $\text{cost}(b)$. $\square$

The blocks are the connected components of a bipartite graph with at most n edges, and the block costs need one pass over E^+ . Hence χ is computed in $O(n + m)$ time. The blocks form the join $a \vee b$ of the two partitions in the partition lattice: the finest partition whose parts are unions of clusters of a and also of b . It is the finest partition for which every block-wise choice between the parents yields each parent as one of the choices and leaves the pairs between blocks separated in both. Decomposing two parents into independent components and taking the better parent in each is the partition crossover of Whitley et al. [70, 71], Tinós et al. [72] for tours and of Tinós et al. [73] for pseudo-Boolean functions; recombination through the overlap of groups is common in grouping genetic algorithms [74, 75], memetic graph partitioning and clustering [76, 77] and ensemble clustering [78], and fusion moves [57] combine proposals optimally within a restricted space. Theorem 12 combines *any* two clusterings, including clusterings produced by different solvers. On our graphs, two independent runs typically differ in thousands of blocks, most of them ties, and each run wins some of the others.

6.3 Annealing moves

The search moves nodes of the contracted instance. Let $S_X = \sum_{K \in X} s_K$ be the size of a cluster X and $w(K, X) = \sum_{L \in X \setminus \{K\}} w_{KL}$. Moving a node K of size s from its cluster A to a cluster B changes cost_w by

$$\delta = s(S_B - S_A + s) + 2(w(K, A) - w(K, B)). \quad (7)$$

Exchanging $K \in A$ with $L \in B$ changes it by the sum of two such terms, where the second one uses $S_A - s_K$, $S_B + s_K$, $w(L, B) + w_{KL}$ and $w(L, A) - w_{KL}$. Both are evaluated in $O(\deg K + \deg L)$ time. The simulated annealing [79] proposes the move to the cluster of a random neighbour, the best move, or a move to a new cluster, and accepts with the Metropolis rule under a geometric cooling schedule. It records the best state with an undo journal, so a new best state costs $O(1)$ instead of a copy of the labelling. This alone made the annealing 7.3 times faster on **com-Amazon**.

Region-wise acceptance.

Theorem 12 turns annealing into an iterated local search with a decomposed acceptance test. From the current clustering x we anneal a copy, starting at a raised temperature, down to its final state y , and set $x \leftarrow \chi(x, y)$. Every block in which y improved on x is kept, even if y is worse overall, and the cost of x never increases.

Localized annealing with rollback.

A step grows a region of at most 30 nodes by breadth-first search around a centre, anneals only the nodes of the region from a high to a low temperature, finishes with

Algorithm 2 PXM_{em}(G^+, T)

```
1: contract the critical cliques
2: for  $i = 1, \dots, 4$  do
3:    $p_i \leftarrow \text{PIVOT}(G^+)$ ;  $f_i \leftarrow$  one round of  $\text{ITERATEDFLIP}(p_i)$ 
4:    $x_i \leftarrow$  the cheaper of  $p_i$  and  $f_i$  on the contracted nodes; anneal  $x_i$ 
5: end for
6: while time remains do
7:   choose an operator by an  $\varepsilon$ -greedy bandit on the recent gain per second
8:   crossover: parents  $a, b$  by binary tournaments;  $z \leftarrow \chi(a, b)$ ;
      contract the blocks common to  $a$  and  $b$  and run a multilevel local search from
       $z$ ;
      anneal on the nodes where  $a$  and  $b$  disagree; accept region-wise; crowding
      replacement
9:   PX-annealing step on a member (temperature chosen by a second bandit)
10:  localized annealing with rollback on a member
11: end while
12: return the cheapest member, expanded to  $V$ 
```

greedy best moves, and undoes all moves if their total change is positive. The undo is exact: moves are reversed in LIFO order together with the free list of cluster labels. Half of the centres are drawn with probability proportional to the degree. Comparing our solutions with those of KaPoCE block by block showed why: the blocks in which KaPoCE was better all contained a vertex of high degree.

6.4 The memetic search

Algorithm 2 keeps a population of four clusterings. A crossover child enters the population only by replacing a member that is not cheaper, preferring the most similar members (crowding); the other two operators replace a member only if they improve it. Contracting the blocks common to both parents and continuing with a multilevel search from the recombined clustering follows the recombination of memetic graph partitioners and clusterers [9, 76, 77]. The flipping round of line 3 costs a few full local-search passes whatever its time limit; with budgets below 300 s the seeds are therefore Pivot clusterings only, and no further seed is generated once half of the budget is used.

Theorem 13 *Let $\mathcal{C}$ be the output of Algorithm 2 and p_1 the first Pivot clustering. Then $\text{cost}(\mathcal{C}) \leq \text{cost}(p_1)$ for every outcome of the random choices, and hence $\mathbb{E}[\text{cost}(\mathcal{C})] \leq 3 \text{OPT}$. Together with a lower bound LB from Section 3.2, $\text{cost}(\mathcal{C})/\lceil \text{LB} \rceil$ bounds the ratio achieved on the instance.*

The guarantee is inherited from the Pivot seed, as for any procedure that starts from a Pivot clustering and never makes its best solution worse; we state it because every operator of the search must respect it, and this is what the formal development checks.

Proof The seed x_1 costs at most $\text{cost}(p_1)$: Pivot does not separate twins, so its contraction preserves its cost, and x_1 is the cheaper of the two candidates. The annealing returns the cheapest state of a trajectory that starts at its input; region-wise acceptance does not increase the cost by Theorem 12; the localized annealing keeps a step only if its total change is non-positive; and every replacement is by a clustering that is not more expensive than the member it replaces. By induction over the replacements, some member always costs at most $\text{cost}(x_1)$, so the cheapest final member costs at most $\text{cost}(p_1)$. Taking expectations and applying the Pivot guarantee [3] gives the second claim. The last claim is Theorem 5(a). $\square$

7 Computational study

7.1 Setup

Implementation and machines.

All methods are implemented in Python with Numba-compiled kernels (package `ccbench`). LPs and MIPs are solved with HiGHS 1.15 [66] and graph partitions are computed with METIS [65]. All runs of this section were made on Google Colab virtual machines with two logical cores and 12.7 GB of memory, one run per machine at a time, with the CPU model recorded in every result file (Intel® Xeon® CPU @ 2.20GHz). Every run uses one thread. Budgets are wall-clock times; CPU times are recorded as well.

Instances.

Table 1 lists the SNAP graphs [80], used as complete instances with E^+ given by the edges (for signed networks, by the positive edges [81, 82]). Fifteen of them were used during development; the twelve held-out graphs were chosen afterwards to span 10^3 to 10^6 vertices and several graph classes, and were fixed before any run on them (Appendix B). We also use the 200 instances of each track of PACE 2021 [8]: the exact track ($n \leq 620$; 173 instances with known optimum and 27 open) and the heuristic track (up to $2 \cdot 10^6$ vertices). Ten heuristic-track instances coincide with development graphs and seven with held-out graphs (Table 1).

Lower bounds compared.

All bounds are rounded up to integers. Bounds marked *checked* are values returned by the checker of Section 4 from an archived certificate.

- *Triangle packings*: a maximal pair-disjoint packing of bad triangles (the bound of MatchFlipPivot [14]) and a fractional packing by the Garg–Könemann scheme with $\varepsilon = 0.05$ [28, 29], scaled by its maximum load.
- *Sparse star packing* (checked): the greedy star packing on P improved by ruin and recreate, installed as a dual solution without any LP. The number of iterations is set before the run from the iteration rate measured on the instance, so as to fill the budget; the measured running times are reported. This is the scalable form of the root star-packing bound of the KaPoCE branch-and-bound.
- *Metric LP on P* : LP_P with triangle rows (1)–(2) only, solved by cutting planes with HiGHS until no row is violated or the budget ends; and the same with star

Table 1 SNAP instances: vertices, edges, size of the distance-two support P and role. “–”: the support exceeds $3 \cdot 10^7$ pairs and does not fit into memory together with the block LPs; no bound is computed. PACE: the heuristic-track instance with the same graph.

graph	n	m	$ P $	role	PACE
ca-GrQc	5,242	14,484	78,224	dev	heur167
ca-HepTh	9,877	25,973	205,421	dev	heur173
ca-HepPh	12,008	118,489	1,638,915	dev	heur175
ca-AstroPh	18,772	198,050	4,639,091	dev	heur176
ca-CondMat	23,133	93,439	1,169,993	dev	heur178
email-Enron	36,692	183,831	15,241,801	dev	heur179
loc-Brightkite	58,228	214,078	8,576,618	dev	
soc-Epinions	75,879	405,740	–	dev	heur180
Slashdot+	75,144	382,882	–	dev	
Epinions+	114,467	592,591	–	dev	
BitcoinOTC+	5,573	18,591	972,669	dev	
BitcoinAlpha+	3,683	12,972	493,107	dev	
com-Amazon	334,863	925,872	6,933,597	dev	heur193
com-DBLP	317,080	1,049,866	13,718,957	dev	heur191
com-Youtube	1,134,890	2,987,624	–	dev	heur198
email-Eu-core	1,005	16,064	223,665	held-out	heur094
facebook	4,039	88,234	1,446,301	held-out	heur166
wiki-Vote	7,115	100,762	3,473,778	held-out	
cit-HepTh	27,770	352,285	–	held-out	
cit-HepPh	34,546	420,877	11,577,385	held-out	
p2p-Gnutella31	62,586	147,892	1,663,263	held-out	
soc-Slashdot0902	82,168	504,230	–	held-out	heur186
loc-Gowalla	196,591	950,327	–	held-out	heur188
email-EuAll	265,214	364,481	–	held-out	
amazon0302	262,111	899,792	6,588,959	held-out	heur189
web-NotreDame	325,729	1,090,108	–	held-out	heur192
roadNet-PA	1,088,092	1,541,898	4,421,718	held-out	heur197

rows. These are LP values in floating point, reported to separate the strength of the relaxation from the dual method.

- *CertiFlip bound* (checked): Theorem 5, as computed within CertiFlip (half of the budget), and, on a subset of PACE instances, with a long budget and one block covering the graph.
- *KaPoCE root bounds*: the star-packing and P_3 -packing bounds of the KaPoCE branch-and-bound as published by its authors [10]. Their dense-matrix implementation does not run on the SNAP graphs beyond about 10^4 vertices.

Solvers compared.

Pivot [3], the best of 50 Pivot runs, MatchFlipPivot [14], vertex-move local search [49], our multilevel Louvain-type local search, Leiden with the constant Potts model at resolution $1/2$ [50, 52], our iterated flipping local search, CertiFlip, PXMem,

and KaPoCE [7], the winner of both PACE 2021 tracks (heuristic version, commit 64e2101). SCC [9] optimises the objective of sparse signed graphs, in which a pair without an edge costs nothing, whereas here every non-edge is a negative pair. We run it (commit 05d0849, memetic version with one process) on the complete encoding, in which every non-adjacent pair is an explicit edge of weight -1 . This needs $\binom{n}{2}$ edges and is feasible on the PACE exact track (Section 7.2) but not on the SNAP graphs.

Head-to-head protocol.

PXMem and KaPoCE are compared under a protocol that was fixed before the runs (`experiments/heldout.md`). On one machine, the two solvers run one after the other in an order drawn from a hash of (graph, seed), each on the same logical core while the machine is otherwise idle. Both read the same PACE-format text file inside their measured time. KaPoCE’s only changes are a seed and its two internal time limits, set to $\frac{590}{600}T$ and $\frac{90}{600}T$ (the ratio of its PACE configuration); it is stopped with SIGTERM at T and its output is validated. PXMem is scored at KaPoCE’s elapsed wall-clock time, read from its complete improvement history, so the comparison is at equal elapsed time; its value at T is also recorded. Invalid KaPoCE output counts as a PXMem win, as in PACE, and is reported separately. Budgets are $T = 600$ s (primary), 150 s and 60 s, with ten seeds on each held-out SNAP graph and one run on each PACE heuristic instance at 600 s.

Statistics.

The unit of inference is the graph. For each graph we take the median over seeds of the relative difference $(\text{cost}_{\text{PXMem}} - \text{cost}_{\text{KaPoCE}}) / \text{cost}_{\text{KaPoCE}}$ and apply the two-sided Wilcoxon signed-rank test [83] and the sign test to the twelve medians. Per graph we report the exact sign test over seeds with Holm’s correction [84] over graphs, noting that with ten seeds only a 10:0 split can be significant, and the Hodges–Lehmann estimate [85] of the paired difference with its exact confidence interval. We do not claim equivalence from non-significance. Performance profiles [86] summarise the PACE heuristic track.

7.2 Lower bounds on the PACE exact track

[pending] Discussion of Tables 2 and 3.

7.3 Certified gaps on SNAP graphs

Certified gaps.

On the 18 graphs whose support fits into memory, the larger of the two checked bounds certifies the archived clusterings within 0.25% to 23.5% (median 6.6%; Table 5). The triangle packings, the lower bounds that scale without the support, only certify factors of 1.17 to 1.90 on the same graphs, which is the regime of the MatchFlipPivot guarantee. Across the three CertiFlip seeds the bound varies by at most 0.99%. The largest gaps are on the held-out social and citation graphs **facebook**, **email-Eu-core**, **wiki-Vote** and **cit-HepPh**; the column gap* shows that no solver we ran closes them from above, so they are due to the lower bound, the upper bound, or both, and we

Table 2 PACE 2021 exact track, the 173 instances with known optimum: lower bounds. Checked bounds are values returned by the checker from archived certificates (budget 60 s, of which CertiFlip spends at most half on the bound). The metric LP on P is reported only where cutting planes converged within 900 s (every fourth instance was run). The last two rows are the root bounds of the KaPoCE branch-and-bound as published by its authors.

bound	instances	mean LB/OPT	min LB/OPT	LB = OPT
greedy triangle packing	173	0.8612	0.660	14
fractional triangle packing	173	0.9664	0.811	30
sparse star packing	[pending]			
CertiFlip bound	173	0.9874	0.816	106
metric LP on P , triangle rows	[pending]			
CertiFlip bound, 1800 s	[pending]			
B&B star packing [10]	173	0.9948	0.891	79
B&B P_3 packing [10]	173	0.9742	0.813	59

Table 3 Mean LB/OPT on the solved PACE exact instances by edge density $m/\binom{n}{2}$.

edge density	instances	CertiFlip	sparse star packing	B&B star packing
[0, 0.1)	35	0.9978	—	0.9934
[0.1, 0.3)	40	0.9781	—	0.9971
[0.3, 0.5)	48	0.9823	—	0.9914
[0.5, 1)	50	0.9923	—	0.9974

Table 4 PACE 2021 exact track (200 instances): primal quality against the best known value (published upper bound or better value of any run). CertiFlip: the certificate runs of Table 2, one seed on the Colab machines; SCC: one seed on the complete encoding, on a 4-core machine with three runs in parallel; ^a: best of three seeds from earlier runs on the same 4-core machine with four runs in parallel.

method	best known	mean gap (%)	max gap (%)
CertiFlip (60 s)	192/200	0.006	0.32
SCC [9] (60 s)	200/200	0.000	0.00
KaPoCE (60 s) ^a	200/200	0.000	0.00
IteratedFlip ^a	187/200	0.007	0.32
Leiden-CPM ^a	158/200	0.119	3.00
Pivot ^a	13/200	32.635	229.54

cannot tell which. The median gap is 6.1% on the development graphs and 9.2% on the held-out graphs. We attribute the difference mainly to the composition of the held-out set, which contains the denser social and citation networks, but with seven held-out graphs we cannot exclude that the method is tuned to the development graphs.

Star packing against the block LPs.

The sparse star packing alone gives a larger checked bound than the CertiFlip bound on 13 of the 18 graphs, although the packing run took a median of 205 s (47 to 1285 s;

Table 5 SNAP graphs whose support fits into memory: checked lower bounds and certified gaps. UB: the best archived clustering (re-checked from its file). CertiFlip: the best checked bound of three CertiFlip runs of 600 s (seeds 0–2), of which the bound takes at most half; packing: the checked bound of one run of the sparse star packing alone (seed 0); the larger of the two is in bold. gap: $(UB - LB)/LB$ with the larger checked bound; gap*: the same with the best value of any run of any solver (reported, not archived); tri. LP: the same with the metric LP on P with triangle rows, where cutting planes converged within 1,800 s (“—” otherwise); tri. factor: best known value divided by the better of the two triangle packings. Upper block: development graphs; lower block: held-out graphs.

graph	UB	checked LB		gap (%)	gap* (%)	tri. LP gap (%)	tri. factor
		CertiFlip	packing				
ca-GrQc	6,051	6,036	6,003	0.25	0.23	22.69	1.28
ca-HepTh	15,861	15,682	15,705	0.99	0.69	—	1.45
ca-HepPh	46,194	43,199	42,865	6.93	6.79	—	1.17
ca-AstroPh	125,682	118,224	117,865	6.31	6.19	—	1.46
ca-CondMat	56,974	55,283	55,627	2.42	2.17	—	1.43
email-Enron	150,701	140,551	140,854	6.99	6.79	—	1.77
loc-Brightkite	175,081	163,366	163,927	6.80	6.57	—	1.71
BitcoinOTC+	16,497	15,582	15,642	5.47	5.24	—	1.80
BitcoinAlpha+	11,461	10,749	10,802	6.10	5.88	—	1.81
com-Amazon	615,035	562,505	574,956	6.97	6.50	—	1.43
com-DBLP	608,183	582,821	589,979	3.09	2.84	—	1.42
email-Eu-core	12,732	11,657	11,413	9.22	9.22	—	1.59
facebook	53,547	43,344	42,953	23.54	23.54	—	1.37
wiki-Vote	95,263	85,487	85,522	11.39	11.39	—	1.90
cit-HepPh	357,463	322,666	324,327	10.22	10.20	—	1.72
p2p-Gnutella31	131,362	130,297	130,336	0.79	0.79	—	1.79
amazon0302	540,844	503,032	511,221	5.79	5.79	—	1.32
roadNet-PA	944,542	865,131	884,166	6.83	6.83	—	1.30

Table 6 The metric LP restricted to P , solved by cutting planes with HiGHS on one core, against the best checked bound of Table 5. LP values are in floating point and not certified; a value in parentheses did not converge within the budget and is not a bound.

graph	best	best checked	LP _P , triangle rows		LP _P , + star rows	
	known	LB	value	time (s)	value	time (s)
ca-GrQc	6,050	6,036	4,931	2	6,039	2578

its number of ruin-and-recreate iterations is fixed in advance from a measured iteration rate) and the CertiFlip bound a median of 289s. The CertiFlip bound is larger only on **ca-GrQc**, **ca-HepPh**, **ca-AstroPh**, **email-Eu-core**, **facebook**, which are small or dense enough for the block LPs to converge. Within CertiFlip the packing phase is short, and on large sparse graphs the block LPs, started from a weaker packing, do not make up the difference in the remaining time. On these graphs the certified gap is therefore due to the packing; the LP part of the method matters on the dense graphs and on the PACE instances (Section 7.2). Running a long packing first and the block

Table 7 PXMem against KaPoCE on the twelve held-out SNAP graphs, $T = 600$ s (primary budget), ten seeds per graph. W/T/L: runs in which PXMem is better, equal, worse. median, HL: median and Hodges–Lehmann estimate of $(\text{cost}_{\text{PXMem}} - \text{cost}_{\text{KaPoCE}}) / \text{cost}_{\text{KaPoCE}}$ in percent, negative when PXMem is better, with the exact 95% confidence interval; p : exact sign test over the seeds; p_{Holm} : Holm-adjusted over the twelve graphs. Last row: the pre-registered tests over the twelve graph medians.

graph	runs	W/T/L	median (%)	HL (%)	[95% CI]	p	p_{Holm}
email-Eu-core	10	0/10/0	+0.000	+0.000	[+0.000, +0.000]	1.000	1.000
facebook	10	0/10/0	+0.000	+0.000	[+0.000, +0.000]	1.000	1.000
wiki-Vote	10	9/0/1	-0.016	-0.015	[-0.023, -0.007]	0.021	0.193
cit-HepTh	10	9/0/1	-0.016	-0.016	[-0.024, -0.008]	0.021	0.193
cit-HepPh	10	9/0/1	-0.019	-0.019	[-0.029, -0.010]	0.021	0.193
p2p-Gnutella31	10	0/7/3	+0.000	+0.000	[+0.000, +0.000]	0.250	1.000
soc-Slashdot0902	10	1/1/8	+0.003	+0.003	[+0.001, +0.005]	0.039	0.234
loc-Gowalla	10	10/0/0	-0.010	-0.011	[-0.016, -0.006]	0.002	0.023
email-EuAll	10	4/0/6	+0.003	+0.002	[-0.001, +0.005]	0.754	1.000
amazon0302	10	0/0/10	+0.001	+0.001	[+0.001, +0.002]	0.002	0.023
web-NotreDame	10	3/0/7	+0.001	+0.001	[-0.002, +0.003]	0.344	1.000
roadNet-PA	10	0/0/10	+0.084	+0.083	[+0.064, +0.099]	0.002	0.023
graph medians: 4 better, 3 equal, 5 worse; Wilcoxon $p = 0.734$, sign test $p = 1.000$							

LPs afterwards is the obvious combination. We did not tune or evaluate it before fixing the held-out protocol and do not report it. The gaps of Table 5 use the better of the two certificates, which were computed in two separate runs.

Strength of the relaxation.

Table 6 separates the strength of the relaxation from the way it is solved. On **ca-GrQc** the metric LP on P with triangle rows converges in 2 s to 4,931, 22.7% below the best known value. With star rows the LP reaches 6,039, 0.18% below the best known value and slightly above the best checked bound (6,036), after 2578 s. The triangle inequalities, which is what the classical metric LP has, are therefore far from sufficient on this graph; the star inequalities close almost all of the gap, and the anytime dual reaches nearly the same value in a tenth of the time.

7.4 PXMem against KaPoCE

Primary budget.

At 600 s (Table 7, 120 paired runs) the graph medians favour PXMem on 4 graphs and KaPoCE on 5, with 3 ties. Neither pre-registered test rejects equality of the two solvers (Wilcoxon $p = 0.734$, sign test $p = 1.000$). After Holm’s correction PXMem is better on **loc-Gowalla** and KaPoCE on **amazon0302** and **roadNet-PA**; PXMem wins nine of ten seeds on **wiki-Vote**, **cit-HepTh** and **cit-HepPh**, which is not significant after the correction. All differences are small: the largest graph median is 0.08%, and on the held-out graphs with a certificate the certified gap of the best solution is at least 0.79% (Table 5). The certificates therefore cannot separate the two solvers, and neither solver is shown to be near-optimal on these graphs beyond what Table 5 certifies. KaPoCE returned a valid solution in every run (0 forfeits). Scoring PXMem

Table 8 As Table 7, $T = 150$ s.

graph	runs	W/T/L	median (%)	HL (%)	[95% CI]	p	p_{Holm}
email-Eu-core	10	0/8/2	+0.000	+0.000	[+0.000, +0.004]	0.500	1.000
facebook	10	1/8/1	+0.000	+0.000	[−0.003, +0.004]	1.000	1.000
wiki-Vote	10	9/0/1	−0.010	−0.012	[−0.019, +0.002]	0.021	0.150
cit-HepTh	10	9/0/1	−0.020	−0.018	[−0.031, +0.008]	0.021	0.150
cit-HepPh	10	7/0/3	−0.010	−0.009	[−0.021, +0.009]	0.344	1.000
p2p-Gnutella31	10	0/5/5	+0.000	+0.000	[+0.000, +0.002]	0.062	0.312
soc-Slashdot0902	10	0/0/10	+0.007	+0.007	[+0.005, +0.009]	0.002	0.023
loc-Gowalla	10	4/0/6	+0.003	+0.004	[−0.001, +0.007]	0.754	1.000
email-EuAll	10	10/0/0	−0.007	−0.008	[−0.017, −0.003]	0.002	0.023
amazon0302	10	0/0/10	+0.007	+0.007	[+0.005, +0.008]	0.002	0.023
web-NotreDame	10	0/0/10	+0.021	+0.022	[+0.017, +0.028]	0.002	0.023
roadNet-PA	10	0/0/10	+0.463	+0.439	[+0.365, +0.517]	0.002	0.023
graph medians: 4 better, 2 equal, 6 worse; Wilcoxon $p = 0.846$, sign test $p = 0.754$							

Table 9 As Table 7, $T = 60$ s.

graph	runs	W/T/L	median (%)	HL (%)	[95% CI]	p	p_{Holm}
email-Eu-core	4	0/1/3	+0.008	−		0.250	1.000
facebook	4	1/2/1	+0.000	−		1.000	1.000
wiki-Vote	4	0/0/4	+0.023	−		0.125	1.000
cit-HepTh	4	0/1/3	+0.003	−		0.250	1.000
cit-HepPh	4	1/0/3	+0.005	−		0.625	1.000
p2p-Gnutella31	4	0/0/4	+0.003	−		0.125	1.000
soc-Slashdot0902	4	0/0/4	+0.020	−		0.125	1.000
loc-Gowalla	4	0/0/4	+0.025	−		0.125	1.000
email-EuAll	4	1/0/3	+0.025	−		0.625	1.000
amazon0302	4	0/0/4	+0.030	−		0.125	1.000
web-NotreDame	4	0/0/4	+0.066	−		0.125	1.000
roadNet-PA	3	0/0/3	+2.641	−		0.250	1.000
graph medians: 0 better, 1 equal, 11 worse; Wilcoxon $p = 0.001$, sign test $p = 0.001$							

at T instead of at KaPoCE’s elapsed time changes the graph medians to 4 better and 5 worse (Wilcoxon $p = 0.734$). We do not claim that the two solvers are equivalent: the study was not designed as an equivalence test.

Shorter budgets.

At 150 s (Table 8) the picture is similar (4 better, 2 equal, 6 worse; Wilcoxon $p = 0.846$), with larger losses of PXM on the largest graphs, up to 0.46% on **roadNet-PA**. At 60 s (Table 9) KaPoCE is better on 11 of the 12 graph medians (Wilcoxon $p = 0.001$). PXM is a population method, and a short budget leaves few recombinations on graphs with 10^5 vertices or more; we did not analyse this further. PXM is thus not a replacement for KaPoCE at short budgets, and at the primary budget the two are not distinguishable by our tests.

PACE heuristic track.

On the 200 instances of the PACE 2021 heuristic track, one run per instance at 600 s under the same protocol, the instance is the unit. The **[pending]** instances that

Table 10
PACE 2021
heuristic
track,
 $T = 600$ s,
one run per
instance:
PXMem
against
KaPoCE by
instance size.
W/T/L:
instances on
which
PXMem is
better,
equal, worse;
 $\sum$: total
difference of
the costs;
median: of
the relative
differences in
percent.

[pending]

[pending]

Fig. 5 Performance profile [86] on the PACE heuristic track without the development twins: fraction of instances on which a solver’s cost is within a factor τ of the better of the two.

coincide with development graphs are reported separately (Table 10). On the other [pending] instances PXMem is better on [pending], equal on [pending] and worse on [pending] (Wilcoxon $p =$ [pending], sign test $p =$ [pending]); KaPoCE returned no valid solution on [pending] instances, which count as PXMem wins as in PACE. Figure 5 shows the performance profile.

8 Conclusion

Lower bounds that scale to large correlation clustering instances used to certify only factors of about two. Three observations change this on sparse graphs: the metric LP and its strengthenings can be restricted to the pairs at distance at most two; dual block-coordinate ascent turns the restricted relaxation into an anytime bound whose state is linear in the support; and star packings on the sparse support are a cheap and strong starting point. The resulting bounds are exported as certificates that an independent checker verifies against the raw instance in exact arithmetic. The gap decomposition turns the same certificate into a map of where a clustering can still improve.

Limitations.

- On dense instances the root star-packing bound of the KaPoCE branch-and-bound remains stronger than ours on average; there the block LPs are too large to converge within the budget.
- On 13 of the 18 SNAP graphs the sparse star packing alone gives a larger bound than CertiFlip’s combination of a short packing and block LPs; there the LP part of the method does not pay for its time within the budget (Section 7.3).
- The support has $O(\sum_v \deg(v)^2)$ pairs. On graphs with strong hubs it does not fit into memory and no bound is computed. Restricting the support to centres of bounded degree would keep the bound valid.
- The checker is short but not formally verified; the formal development covers the inequality it evaluates, not its code (Appendix A).
- The a-priori guarantees of our primal methods are inherited from Pivot, and the 2.06 guarantee applies only when one block covers the graph and separation converges.

Open directions.

Parallel block solves; cluster-LP columns inside blocks, to which the roundings of Cao et al. [5] and García-Soriano and Schohn [6] could be applied; overlapping blocks with Lagrangian cost splitting; and a verified checker, for example by extracting it from the Lean development.

Supplementary information. The code, the instance lists, the raw results of every run and the archived certificates are available at https://github.com/vinhqdang/correlation_clustering; the file REPRODUCE.md lists one command per table and figure.

Declarations

- Funding: not applicable.
- Competing interests: the author declares no competing interests.
- Data availability: all instances are public (PACE 2021 and SNAP); the repository contains scripts that download them and check their SHA-256 hashes.
- Code availability: see the supplementary information.

Appendix A Formal development

The correctness of the solvers rests on short combinatorial arguments and on exact formulas that the implementation evaluates millions of times per second. An error in a sign or a boundary case would not crash the code; it would silently produce worse or invalid results. We formalised part of the mathematics in Lean 4 [87] with Mathlib [88]. The development (`lean/CCProofs`, about 1,800 lines in fourteen files) is checked by the Lean kernel without `sorry` and depends only on the axioms `propext`, `Classical.choice` and `Quot.sound`. Table A1 lists what it covers. It does *not* cover Theorem 3, the block step and the convergence statement of Theorem 5(b,c), the gap decomposition (Theorem 7) or its local corollary, the 2.06 part of Theorem 9, or any code. The Pivot guarantee enters as a hypothesis.

The formal model is that of Section 2: a finite vertex set, a simple graph G^+ and a labelling $c : V \rightarrow \alpha$ into a type with decidable equality. The cost is summed over ordered pairs, so every count appears multiplied by two (`cost_eq_two_mul_card` relates it to the unordered count). The weighted instance has integer node sizes and symmetric weights, as in (6).

The formal statements connect to the code in two places. The bounds we report are values of the expression that `far_dual_bound` bounds, evaluated by the checker of Section 4. The move, swap and contraction formulas are tested against recomputation: exact costs of random clusterings on contracted instances, and the tracked cost against a full recomputation at the end of annealing runs with many moves, swaps and rollbacks. These tests check final states, not every step.

Appendix B Development log

This appendix records which data informed which decision, so that the reader can judge which results are out of sample.

- *Development data.* The fifteen SNAP graphs of Table 1 marked “dev” and the 200 PACE 2021 exact-track instances. All parameters of the bound (block size 2500, $|H| \leq 8$ for local subgraph rows, ruin regions of about 12 vertices, half of the budget for the bound), of the LNS ($|B| \leq 60$) and of PXMmem (population 4, regions of at most 30 nodes, bandit settings, the degree bias of the region centres) were chosen by looking at runs on these instances. Several PXMmem components were added after inspecting where KaPoCE’s solutions were better than ours on development graphs (Section 6).
- *Earlier comparisons.* A head-to-head comparison with KaPoCE on the development graphs and on the PACE heuristic track was carried out before the held-out set was fixed. Its raw results remain in the repository; it is not the basis of any claim in this paper, because the same graphs were used for tuning and because it measured the two solvers asymmetrically (KaPoCE’s time included reading the input, PXMmem’s did not, and the order was fixed).
- *Held-out data.* The twelve SNAP graphs marked “held-out” in Table 1 and the analysis plan of Section 7.1 were fixed and committed to the repository (`experiments/heldout.md`) before any run of either solver on them. No parameter was changed afterwards. The PACE heuristic-track instances with the same vertex count as a development graph (ten instances) are excluded from the held-out PACE summary.

References

- [1] Bansal, N., Blum, A., Chawla, S.: Correlation clustering. Machine Learning **56**(1-3), 89–113 (2004) <https://doi.org/10.1023/B:MACH.0000033116.57574.95> . Conference version: FOCS 2002, pp. 238–247, doi:10.1109/SFCS.2002.1181947
- [2] Charikar, M., Guruswami, V., Wirth, A.: Clustering with qualitative information. Journal of Computer and System Sciences **71**(3), 360–383 (2005) <https://doi.org/10.1016/j.jcss.2005.06.001>

Table A1 Coverage of the formal development (Lean 4, Mathlib). “Full”: the statement of the paper is proved as stated (up to the ordered-pair convention); “partial”: the part named is proved; “reduction”: proved from the cited guarantee taken as a hypothesis; “–”: not formalised.

result	coverage	Lean theorems
Lemma 1	partial	far pairs are separated by every optimal clustering: <code>optimal_separates_far</code> , <code>sepFar_of_optimal</code>
Theorem 3	–	
Proposition 4(i)	full	<code>star_row_valid</code> (star rows in x -form, far-separating clusterings)
Proposition 4(ii)	partial	bad triangles and induced stars: <code>bad_triangle_costIn</code> , <code>star_bound</code> ; packings: <code>packing_bound</code>
rows (1)–(2)	full	<code>sep_triangle</code> , <code>sep_far_triangle</code>
Theorem 5(a)	full	<code>far_dual_bound</code> , <code>far_dual_le_opt</code> (bound on P for rows valid for far-separating clusterings, $\leq$ the cost of every clustering); <code>cost_eq_supp</code>
Theorem 5(b,c), Theorem 7, Corollary 8	–	
Theorem 9	(a) reduction, (b) full, (c) –	(a) as Theorem 13; (b) is Theorem 5(a)
Proposition 6	partial	the inequality evaluated by the checker: <code>far_dual_bound</code> , <code>far_dual_le_opt</code> ; the checker code itself is not verified
Lemma 11	full	<code>twins_split_improvable</code> , <code>optimal_keeps_twins</code>
Pivot keeps twins together	partial	one round: <code>pivot_round_twins</code>
contraction (6)	partial	cost identity: <code>cost_contract</code>
move and swap formulas (7)	full	<code>cost_move</code> , <code>costW_move</code> , <code>costW_swap</code> , <code>cost_update</code>
Theorem 12	full	any pairwise-additive cost: <code>crossoverP_le</code> ; unweighted: <code>crossover_le</code> ; weighted: <code>crossoverW_le</code> ; relabelling: <code>cost_comp_injective</code> , <code>costW_comp_injective</code>
Theorem 13	reduction	<code>reach_le</code> , <code>output_le</code> , <code>search_guarantee</code> : non-worsening replacements keep a member no worse than the seed; the Pivot guarantee [3] is a hypothesis

[org/10.1016/j.jcss.2004.10.012](https://doi.org/10.1016/j.jcss.2004.10.012) . Conference version: FOCS 2003, pp. 524–533, doi:10.1109/SFCS.2003.1238225

- [3] Ailon, N., Charikar, M., Newman, A.: Aggregating inconsistent information: Ranking and clustering. *Journal of the ACM* **55**(5), 23–12327 (2008) <https://doi.org/10.1145/1391788.1391792>

doi.org/10.1145/1411509.1411513 . Conference version: STOC 2005, pp. 684–693, doi:10.1145/1060590.1060692

- [4] Chawla, S., Makarychev, K., Schramm, T., Yaroslavtsev, G.: Near optimal LP rounding algorithm for correlation clustering on complete and complete k-partite graphs. In: Proceedings of the 47th Annual ACM Symposium on Theory of Computing (STOC), pp. 219–228 (2015). <https://doi.org/10.1145/2746539.2746604>
- [5] Cao, N., Cohen-Addad, V., Lee, E., Li, S., Newman, A., Vogl, L.: Understanding the cluster linear program for correlation clustering. In: Proceedings of the 56th Annual ACM Symposium on Theory of Computing (STOC), pp. 1605–1616 (2024). <https://doi.org/10.1145/3618260.3649749> . Conference version claims 1.437; arXiv v3 (2025) reports a gap in that proof and proves 1.485
- [6] García-Soriano, D., Schohn, A.: Approximate Dual Separation for the Cluster LP: a 1.387 approximation for Correlation Clustering. arXiv preprint arXiv:2607.27829 (2026)
- [7] Bläsius, T., Fischbeck, P., Gottesbüren, L., Hamann, M., Heuer, T., Spinner, J., Weyand, C., Wilhelm, M.: PACE solver description: KaPoCE: A heuristic cluster editing algorithm. In: 16th International Symposium on Parameterized and Exact Computation (IPEC 2021). LIPIcs, vol. 214, pp. 31–1314 (2021). <https://doi.org/10.4230/LIPIcs.IPEC.2021.31>
- [8] Kellerhals, L., Koana, T., Nichterlein, A., Zschoche, P.: The PACE 2021 parameterized algorithms and computational experiments challenge: Cluster editing. In: Golovach, P.A., Zehavi, M. (eds.) 16th International Symposium on Parameterized and Exact Computation (IPEC 2021). LIPIcs, vol. 214, pp. 26–12618. Schloss Dagstuhl – Leibniz-Zentrum für Informatik, ??? (2021). <https://doi.org/10.4230/LIPIcs.IPEC.2021.26>
- [9] Hausberger, F., Faraj, M.F., Schulz, C.: Scalable multilevel and memetic signed graph clustering. In: Proceedings of the Symposium on Algorithm Engineering and Experiments (ALENEX), pp. 81–94 (2025). <https://doi.org/10.1137/1.9781611978339.7>
- [10] Bläsius, T., Fischbeck, P., Gottesbüren, L., Hamann, M., Heuer, T., Spinner, J., Weyand, C., Wilhelm, M.: A branch-and-bound algorithm for cluster editing. In: 20th International Symposium on Experimental Algorithms (SEA 2022). LIPIcs, vol. 233, pp. 13–11319 (2022). <https://doi.org/10.4230/LIPIcs.SEA.2022.13>
- [11] Veldt, N., Gleich, D.F., Wirth, A., Saunderson, J.: Metric-constrained optimization for graph clustering algorithms. SIAM Journal on Mathematics of Data Science 1(2), 333–355 (2019) <https://doi.org/10.1137/18M1217152> arXiv:1806.01678

- [12] Ruggles, C., Veldt, N., Gleich, D.F.: A parallel projection method for metric constrained optimization. In: Proceedings of the SIAM Workshop on Combinatorial Scientific Computing (CSC), pp. 43–53 (2020). <https://doi.org/10.1137/1.9781611976229.5>
- [13] Sonthalia, R., Gilbert, A.C.: Project and forget: Solving large-scale metric constrained problems. *Journal of Machine Learning Research* **23**(326), 1–54 (2022)
- [14] Veldt, N.: Correlation clustering via strong triadic closure labeling: Fast approximation algorithms and practical lower bounds. In: Proceedings of the 39th International Conference on Machine Learning (ICML). PMLR, vol. 162, pp. 22060–22083 (2022)
- [15] Cohen-Addad, V., Lolck, D.R., Pilipczuk, M., Thorup, M., Yan, S., Zhang, H.: Combinatorial correlation clustering. In: Proceedings of the 56th Annual ACM Symposium on Theory of Computing (STOC), pp. 1617–1628 (2024). <https://doi.org/10.1145/3618260.3649712>
- [16] McConnell, R.M., Mehlhorn, K., Näher, S., Schweitzer, P.: Certifying algorithms. *Computer Science Review* **5**(2), 119–161 (2011) <https://doi.org/10.1016/j.cosrev.2010.09.009>
- [17] Neumaier, A., Shcherbina, O.: Safe bounds in linear and mixed-integer linear programming. *Mathematical Programming* **99**(2), 283–296 (2004) <https://doi.org/10.1007/s10107-003-0433-3>
- [18] Steffy, D.E., Wolter, K.: Valid linear programming bounds for exact mixed-integer programming. *INFORMS Journal on Computing* **25**(2), 271–284 (2013) <https://doi.org/10.1287/ijoc.1120.0501>
- [19] Cheung, K.K.H., Gleixner, A., Steffy, D.E.: Verifying integer programming results. In: Integer Programming and Combinatorial Optimization (IPCO 2017). Lecture Notes in Computer Science, vol. 10328, pp. 148–160. Springer, ??? (2017). https://doi.org/10.1007/978-3-319-59250-3_13
- [20] Wetzler, N., Heule, M.J.H., Hunt, W.A.: DRAT-trim: Efficient checking and trimming using expressive clausal proofs. In: Theory and Applications of Satisfiability Testing – SAT 2014. Lecture Notes in Computer Science, vol. 8561, pp. 422–429. Springer, ??? (2014). https://doi.org/10.1007/978-3-319-09284-3_31
- [21] Gocht, S., Nordström, J.: Certifying parity reasoning efficiently using pseudo-Boolean proofs. In: Proceedings of the AAAI Conference on Artificial Intelligence, vol. 35, pp. 3768–3777. AAAI Press, ??? (2021). <https://doi.org/10.1609/aaai.v35i5.16494>

- [22] Cohen-Addad, V., Lee, E., Newman, A.: Correlation clustering with Sherali-Adams. In: 63rd IEEE Annual Symposium on Foundations of Computer Science (FOCS), pp. 651–661 (2022). <https://doi.org/10.1109/FOCS54457.2022.00068>
- [23] Cohen-Addad, V., Lee, E., Li, S., Newman, A.: Handling correlated rounding error via preclustering: A 1.73-approximation for correlation clustering. In: 64th IEEE Annual Symposium on Foundations of Computer Science (FOCS), pp. 1082–1104 (2023). <https://doi.org/10.1109/FOCS57990.2023.00065>
- [24] Cao, N., Cohen-Addad, V., Li, S., Lee, E., Lolck, D.R., Newman, A., Thorup, M., Vogl, L., Yan, S., Zhang, H.: Solving the correlation cluster LP in sublinear time. In: Proceedings of the 57th Annual ACM Symposium on Theory of Computing (STOC), pp. 1154–1165 (2025). <https://doi.org/10.1145/3717823.3718181>
- [25] Cohen-Addad, V., Lattanzi, S., Mitrović, S., Norouzi-Fard, A., Parotsidis, N., Tarnawski, J.: Correlation clustering in constant many parallel rounds. In: Proceedings of the 38th International Conference on Machine Learning (ICML). PMLR, vol. 139, pp. 2069–2078 (2021)
- [26] Behnezhad, S., Charikar, M., Ma, W., Tan, L.-Y.: Almost 3-approximate correlation clustering in constant rounds. In: IEEE 63rd Annual Symposium on Foundations of Computer Science (FOCS), pp. 720–731 (2022). <https://doi.org/10.1109/FOCS54457.2022.00074>
- [27] Demaine, E.D., Emanuel, D., Fiat, A., Immorlica, N.: Correlation clustering in general weighted graphs. *Theoretical Computer Science* **361**(2-3), 172–187 (2006) <https://doi.org/10.1016/j.tcs.2006.05.008>
- [28] Garg, N., Könemann, J.: Faster and simpler algorithms for multicommodity flow and other fractional packing problems. *SIAM Journal on Computing* **37**(2), 630–652 (2007) <https://doi.org/10.1137/S0097539704446232>
- [29] Fleischer, L.K.: Approximating fractional multicommodity flow independent of the number of commodities. *SIAM Journal on Discrete Mathematics* **13**(4), 505–520 (2000) <https://doi.org/10.1137/S0895480199355754>
- [30] Böcker, S., Briesemeister, S., Klau, G.W.: Exact algorithms for cluster editing: Evaluation and experiments. *Algorithmica* **60**(2), 316–334 (2011) <https://doi.org/10.1007/s00453-009-9339-7>
- [31] Bastos, L., Ochi, L.S., Protti, F., Subramanian, A., Martins, I.C., Pinheiro, R.G.S.: Efficient algorithms for cluster editing. *Journal of Combinatorial Optimization* **31**(1), 347–371 (2016) <https://doi.org/10.1007/s10878-014-9756-7>
- [32] Grötschel, M., Wakabayashi, Y.: A cutting plane algorithm for a clustering problem. *Mathematical Programming* **45**(1-3), 59–96 (1989) <https://doi.org/10.1007/BF01589097>

- [33] Grötschel, M., Wakabayashi, Y.: Facets of the clique partitioning polytope. *Mathematical Programming* **47**(1-3), 367–387 (1990) <https://doi.org/10.1007/BF01580870>
- [34] Miyauchi, A., Sukegawa, N.: Redundant constraints in the standard formulation for the clique partitioning problem. *Optimization Letters* **9**(1), 199–207 (2015) <https://doi.org/10.1007/s11590-014-0754-6>
- [35] Miyauchi, A., Sonobe, T., Sukegawa, N.: Exact clustering via integer programming and maximum satisfiability. In: *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 32 (2018). <https://doi.org/10.1609/aaai.v32i1.11519>
- [36] Koshimura, M., Watanabe, E., Sakurai, Y., Yokoo, M.: Concise integer linear programming formulation for clique partitioning problems. *Constraints* **27**(1-2), 99–115 (2022) <https://doi.org/10.1007/s10601-022-09326-z>
- [37] Irmay, J., Naumann, L.F., Andres, B.: Chorded cycle facets of the clique partitioning polytope. *arXiv preprint arXiv:2411.03407* (2024)
- [38] Sukegawa, N., Yamamoto, Y., Zhang, L.: Lagrangian relaxation and pegging test for the clique partitioning problem. *Advances in Data Analysis and Classification* **7**(4), 363–391 (2013) <https://doi.org/10.1007/s11634-013-0135-5>
- [39] Belyi, A., Sobolevsky, S., Kurbatski, A., Ratti, C.: Subnetwork constraints for tighter upper bounds and exact solution of the clique partitioning problem. *Mathematical Methods of Operations Research* **98**(2), 269–297 (2023) <https://doi.org/10.1007/s00186-023-00835-y> [arXiv:2110.05627](https://arxiv.org/abs/2110.05627)
- [40] Berg, J., Jarvisalo, M.: Cost-optimal constrained correlation clustering via weighted partial Maximum Satisfiability. *Artificial Intelligence* **244**, 110–142 (2017) <https://doi.org/10.1016/j.artint.2015.07.001>
- [41] Guo, J.: A more effective linear kernelization for cluster editing. *Theoretical Computer Science* **410**(8-10), 718–726 (2009) <https://doi.org/10.1016/j.tcs.2008.10.021>
- [42] Cao, Y., Chen, J.: Cluster editing: Kernelization based on edge cuts. *Algorithmica* **64**(1), 152–169 (2012) <https://doi.org/10.1007/s00453-011-9595-1> [arXiv:1008.4250](https://arxiv.org/abs/1008.4250)
- [43] Swoboda, P., Andres, B.: A message passing algorithm for the minimum cost multicut problem. In: *IEEE Conference on Computer Vision and Pattern Recognition (CVPR 2017)*, pp. 4990–4999 (2017). <https://doi.org/10.1109/CVPR.2017.530>
- [44] Abbas, A., Swoboda, P.: RAMA: A rapid multicut algorithm on GPU. In: *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR 2022)*, pp. 8183–8192 (2022). <https://doi.org/10.1109/CVPR52688.2022.00802>

- [45] Abbas, A., Swoboda, P.: ClusterFuG: Clustering fully connected graphs by multicut. In: Proceedings of the 40th International Conference on Machine Learning (ICML). PMLR, vol. 202, pp. 19–30 (2023)
- [46] Lange, J.-H., Karrenbauer, A., Andres, B.: Partial optimality and fast lower bounds for weighted correlation clustering. In: Proceedings of the 35th International Conference on Machine Learning (ICML). PMLR, vol. 80, pp. 2892–2901 (2018)
- [47] Lange, J.-H., Andres, B., Swoboda, P.: Combinatorial persistency criteria for multicut and max-cut. In: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pp. 6086–6095. IEEE, ??? (2019). <https://doi.org/10.1109/CVPR.2019.00625>
- [48] Tourani, S., Shekhovtsov, A., Rother, C., Savchynskyy, B.: MPLP++: Fast, parallel dual block-coordinate ascent for dense graphical models. In: Computer Vision – ECCV 2018. Lecture Notes in Computer Science, vol. 11208, pp. 264–281. Springer, ??? (2018). https://doi.org/10.1007/978-3-030-01225-0_16
- [49] Elsner, M., Schudy, W.: Bounding and comparing methods for correlation clustering beyond ILP. In: Proceedings of the Workshop on Integer Linear Programming for Natural Language Processing, pp. 19–27. Association for Computational Linguistics, ??? (2009). <https://doi.org/10.3115/1611638.1611641>
- [50] Veldt, N., Gleich, D.F., Wirth, A.: A correlation clustering framework for community detection. In: Proceedings of the 2018 World Wide Web Conference (WWW), pp. 439–448 (2018). <https://doi.org/10.1145/3178876.3186110>
- [51] Blondel, V.D., Guillaume, J.-L., Lambiotte, R., Lefebvre, E.: Fast unfolding of communities in large networks. *Journal of Statistical Mechanics: Theory and Experiment* **2008**(10), 10008 (2008) <https://doi.org/10.1088/1742-5468/2008/10/P10008>
- [52] Traag, V.A., Waltman, L., Eck, N.J.: From Louvain to Leiden: guaranteeing well-connected communities. *Scientific Reports* **9**, 5233 (2019) <https://doi.org/10.1038/s41598-019-41695-z>
- [53] Brusco, M.J., Köhn, H.-F.: Clustering qualitative data based on binary equivalence relations: Neighborhood search heuristics for the clique partitioning problem. *Psychometrika* **74**(4), 685–703 (2009) <https://doi.org/10.1007/s11336-009-9126-z>
- [54] Zhou, Y., Hao, J.-K., Goëffon, A.: A three-phased local search approach for the clique partitioning problem. *Journal of Combinatorial Optimization* **32**(2), 469–491 (2016) <https://doi.org/10.1007/s10878-015-9964-9>
- [55] Lu, Z., Zhou, Y., Hao, J.-K.: A hybrid evolutionary algorithm for the clique

- partitioning problem. *IEEE Transactions on Cybernetics* **52**(9), 9391–9403 (2022) <https://doi.org/10.1109/TCYB.2021.3051243>
- [56] Levinkov, E., Kirillov, A., Andres, B.: A comparative study of local search algorithms for correlation clustering. In: *German Conference on Pattern Recognition (GCPR)*. LNCS, vol. 10496, pp. 103–114 (2017). https://doi.org/10.1007/978-3-319-66709-6_9
 - [57] Beier, T., Hamprecht, F.A., Kappes, J.H.: Fusion moves for correlation clustering. In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 3507–3516. IEEE, ??? (2015). <https://doi.org/10.1109/CVPR.2015.7298973>
 - [58] Protti, F., Dantas da Silva, M., Szwarcfiter, J.L.: Applying modular decomposition to parameterized cluster editing problems. *Theory of Computing Systems* **44**(1), 91–104 (2009) <https://doi.org/10.1007/s00224-007-9032-7>
 - [59] Böcker, S., Briesemeister, S., Bui, Q.B.A., Truss, A.: Going weighted: Parameterized algorithms for cluster editing. *Theoretical Computer Science* **410**(52), 5467–5480 (2009) <https://doi.org/10.1016/j.tcs.2009.05.006> . Conference version: COCOA 2008, pp. 1–12, doi:10.1007/978-3-540-85097-7_1
 - [60] Böcker, S.: A golden ratio parameterized algorithm for cluster editing. *Journal of Discrete Algorithms* **16**, 79–89 (2012) <https://doi.org/10.1016/j.jda.2012.04.005> . Conference version: IWOCA 2011, doi:10.1007/978-3-642-25011-8_7
 - [61] Chen, J., Meng, J.: A 2k kernel for the cluster editing problem. *Journal of Computer and System Sciences* **78**(1), 211–220 (2012) <https://doi.org/10.1016/j.jcss.2011.04.001> . Conference version: COCOON 2010, doi:10.1007/978-3-642-14031-0_49
 - [62] Wittkop, T., Baumbach, J., Lobo, F.P., Rahmann, S.: Large scale clustering of protein sequences with FORCE – A layout based heuristic for weighted cluster editing. *BMC Bioinformatics* **8**, 396 (2007) <https://doi.org/10.1186/1471-2105-8-396>
 - [63] Wittkop, T., Emig, D., Lange, S., Rahmann, S., Albrecht, M., Morris, J.H., Böcker, S., Stoye, J., Baumbach, J.: Partitioning biological data with transitivity clustering. *Nature Methods* **7**(6), 419–420 (2010) <https://doi.org/10.1038/nmeth0610-419>
 - [64] Applegate, D., Bixby, R., Chvátal, V., Cook, W.: TSP cuts which do not conform to the template paradigm. In: Jünger, M., Naddef, D. (eds.) *Computational Combinatorial Optimization: Optimal or Provably Near-Optimal Solutions*. LNCS, vol. 2241, pp. 261–303. Springer, ??? (2001). https://doi.org/10.1007/3-540-45586-8_7

- [65] Karypis, G., Kumar, V.: A fast and high quality multilevel scheme for partitioning irregular graphs. *SIAM Journal on Scientific Computing* **20**(1), 359–392 (1998) <https://doi.org/10.1137/S1064827595287997>
- [66] Huangfu, Q., Hall, J.A.J.: Parallelizing the dual revised simplex method. *Mathematical Programming Computation* **10**(1), 119–142 (2018) <https://doi.org/10.1007/s12532-017-0130-5>
- [67] Shaw, P.: Using constraint programming and local search methods to solve vehicle routing problems. In: *Principles and Practice of Constraint Programming – CP98*. LNCS, vol. 1520, pp. 417–431 (1998). https://doi.org/10.1007/3-540-49481-2_30
- [68] Danna, E., Rothberg, E., Le Pape, C.: Exploring relaxation induced neighborhoods to improve MIP solutions. *Mathematical Programming* **102**(1), 71–90 (2005) <https://doi.org/10.1007/s10107-004-0518-7>
- [69] Fischetti, M., Lodi, A.: Local branching. *Mathematical Programming* **98**(1-3), 23–47 (2003) <https://doi.org/10.1007/s10107-003-0395-5>
- [70] Whitley, D., Hains, D., Howe, A.: Tunneling between optima: Partition crossover for the traveling salesman problem. In: *Proceedings of the Genetic and Evolutionary Computation Conference (GECCO)*, pp. 915–922 (2009). <https://doi.org/10.1145/1569901.1570026>
- [71] Whitley, D., Hains, D., Howe, A.: A hybrid genetic algorithm for the traveling salesman problem using generalized partition crossover. In: *Parallel Problem Solving from Nature – PPSN XI*. LNCS, vol. 6238, pp. 566–575. Springer, ??? (2010). https://doi.org/10.1007/978-3-642-15844-5_57
- [72] Tinós, R., Whitley, D., Ochoa, G.: A new generalized partition crossover for the traveling salesman problem: Tunneling between local optima. *Evolutionary Computation* **28**(2), 255–288 (2020) https://doi.org/10.1162/evco_a.00254
- [73] Tinós, R., Whitley, D., Chicano, F.: Partition crossover for pseudo-boolean optimization. In: *Proceedings of the 2015 ACM Conference on Foundations of Genetic Algorithms XIII (FOGA)*, pp. 137–149. ACM, ??? (2015). <https://doi.org/10.1145/2725494.2725497>
- [74] Falkenauer, E.: *Genetic Algorithms and Grouping Problems*. John Wiley & Sons, Chichester, UK (1998)
- [75] Galinier, P., Hao, J.-K.: Hybrid evolutionary algorithms for graph coloring. *Journal of Combinatorial Optimization* **3**(4), 379–397 (1999) <https://doi.org/10.1023/A:1009823419804>
- [76] Sanders, P., Schulz, C.: Distributed evolutionary graph partitioning. In:

- Proceedings of the 14th Workshop on Algorithm Engineering and Experiments (ALENEX), pp. 16–29. SIAM, ??? (2012). <https://doi.org/10.1137/1.9781611972924.2>
- [77] Biedermann, S., Henzinger, M., Schulz, C., Schuster, B.: Memetic graph clustering. In: 17th International Symposium on Experimental Algorithms (SEA 2018). LIPIcs, vol. 103, pp. 3–1315. Schloss Dagstuhl – Leibniz-Zentrum für Informatik, ??? (2018). <https://doi.org/10.4230/LIPIcs.SEA.2018.3>
 - [78] Ovelgönne, M., Geyer-Schulz, A.: An ensemble learning strategy for graph clustering. In: Bader, D.A., Meyerhenke, H., Sanders, P., Wagner, D. (eds.) Graph Partitioning and Graph Clustering. Contemporary Mathematics, vol. 588, pp. 187–205. American Mathematical Society, ??? (2013). <https://doi.org/10.1090/conm/588/11701>
 - [79] Kirkpatrick, S., Gelatt, C.D., Vecchi, M.P.: Optimization by simulated annealing. *Science* **220**(4598), 671–680 (1983) <https://doi.org/10.1126/science.220.4598.671>
 - [80] Leskovec, J., Krevl, A.: SNAP Datasets: Stanford Large Network Dataset Collection. <http://snap.stanford.edu/data> (2014)
 - [81] Leskovec, J., Huttenlocher, D., Kleinberg, J.: Signed networks in social media. In: Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI), pp. 1361–1370 (2010). <https://doi.org/10.1145/1753326.1753532>
 - [82] Kumar, S., Spezzano, F., Subrahmanian, V.S., Faloutsos, C.: Edge weight prediction in weighted signed networks. In: IEEE 16th International Conference on Data Mining (ICDM), pp. 221–230 (2016). <https://doi.org/10.1109/ICDM.2016.0033>
 - [83] Wilcoxon, F.: Individual comparisons by ranking methods. *Biometrics Bulletin* **1**(6), 80–83 (1945) <https://doi.org/10.2307/3001968>
 - [84] Holm, S.: A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics* **6**(2), 65–70 (1979)
 - [85] Hodges, J.L., Lehmann, E.L.: Estimates of location based on rank tests. *The Annals of Mathematical Statistics* **34**(2), 598–611 (1963) <https://doi.org/10.1214/aoms/1177704172>
 - [86] Dolan, E.D., Moré, J.J.: Benchmarking optimization software with performance profiles. *Mathematical Programming* **91**(2), 201–213 (2002) <https://doi.org/10.1007/s101070100263>
 - [87] Moura, L., Ullrich, S.: The Lean 4 theorem prover and programming language. In: Automated Deduction – CADE 28. Lecture Notes in Computer Science, vol. 12699, pp. 625–635. Springer, ??? (2021). <https://doi.org/10.1007/978-3-030-79876-5.37>

- [88] The mathlib Community: The Lean mathematical library. In: Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs (CPP), pp. 367–381 (2020). <https://doi.org/10.1145/3372885.3373824>